CS246—Assignment 1 (Spring 2017)

Due Date 1: Monday, May 15, 5pm
Due Date 2: Friday, May 19, 5pm

Questions 1, 2 and 3 are due on Due Date 1; the remainder of the assignment is due on Due Date 2.

1. Provide a Unix command line to accomplish each of the following tasks. Your answer in each subquestion should consist of a single command or pipeline of commands, with no separating semicolons (;). (Please verify before submitting that your solution consists of a single line. Use `wc` for this.) Before beginning this question, familiarize yourself with the Unix commands outlined on the Unix handout. Keep in mind that some commands have options not listed on the sheet, so you may need to examine some man pages. Note that some tasks refer to a file `myfile.txt`. No `myfile.txt` is given. You should create your own for testing.

   (a) Print the 10th through 25th words (including the 10th and 25th words) in `/usr/share/dict/words`. You may take advantage of the fact that the words in this file are each on a separate line. Place your command pipeline in the file `a1q1a.txt`.

   (b) Print the (non-hidden) contents of the current directory in reverse order. Place your command pipeline in the file `a1q1b.txt`.

   (c) Print the number of lines in the text file `myfile.txt` that do *not* contain the string `cs246`. Place your command pipeline in the file `a1q1c.txt`.

   (d) Print the first line that contains the string `cs246` from the text file `myfile.txt`. Place your command pipeline in the file `a1q1d.txt`.

   (e) Print the number of lines in the text file `myfile.txt` that contain the string `linux.student.cs.uwaterloo.ca` where each letter could be either uppercase or lowercase. Place your command pipeline in the file `a1q1e.txt`.

   (f) Print all (non-hidden) files in any *subdirectory* of the current directory that end with `.c` (immediate subdirectories only, not subdirectories of subdirectories). Do not use `find`. Place your command pipeline in the file `a1q1f.txt`.

   (g) Out of the first 20 lines of `myfile.txt`, how many contain at least one digit? Place the command pipeline that prints this number in the file `a1q1g.txt`.

   (h) Print all (non-hidden) files in the current directory that start with `a`, contain at least one `b`, and end with `.c`. Place your command pipeline in the file `a1q1h.txt`.

   (i) Print a listing, in long form, of all non-hidden entries (files, directories, etc.) in the current directory that are executable by at least one of owner, group, other (the other permission bits could be anything). Do not attempt to solve this problem with `find`. Place your command pipeline in the file `a1q1i.txt`.

2. For each of the following text search criteria, provide a regular expression that matches the criterion, suitable for use with `egrep`. Your answer in each case should be a text file that contains just the regular expression, on a single line (again, use `wc` to verify this). If your pattern contains special characters, enclose it in quotes.

   (a) Lines that contain both cs246 and cs247.
       Place your answer in the file `a1q2a.txt`.

   (b) Lines that contain nothing but a single occurrence of laughter, where laughter is defined as a string of the form `Hahahahahahahahahahahahaha!`, with arbitrarily many `ha`'s.
       Place your answer in the file `a1q2b.txt`.

   (c) Lines that contain nothing but a single occurrence of generalized laughter, which is like ordinary laughter, except that there can be arbitrarily many (but at least one) `a`'s between each pair of consecutive `h`'s. (For example: `Haahahaaaa!`) Place your answer in the file `a1q2c.txt`.

   (d) Lines that contain at least one `a` and at least two `b`'s.
       Place your answer in the file `a1q2d.txt`.

   (e) Lines consisting of a definition of a single C variable of type `int`, without initialization, optionally preceded by `unsigned`, and optionally followed by any single line `//` comment. Example:

       ```
       int  varname;  //  comment
       ```

       You may assume that all of the whitespace in the line consists of space characters (no tabs). You may also assume that `varname` will not be a C keyword (i.e., you do not have to try to check for this with your regular expression). Place your answer in the file `a1q2e.txt`.

3. Write a Bash script called `findGreater` that is executed as: `./findGreater myword file1 file2` where `myword` is a sequence of non-whitespace characters, and `file1` and `file2` are names of text files in the current directory. The script prints the name of the file that contains a larger <u>number of lines</u> with the occurrence of the word `myword`. If the files have the same number of lines that contain occurrences of the provided word, the script prints `file1` followed by a single space and then `file2`. In all cases, the script produces a single line of output to standard output. You may assume that the user will call this script correctly; no error checking is needed. Using the provided `file1.txt` and `file2.txt`, the following shows some example runs of the script (lines not starting with a $ are the output produced by the command executed in the previous line):

   ```
   $ ./findGreater thousand file1.txt file2.txt
   file2.txt
   $ ./findGreater be file1.txt file2.txt
   file1.txt file2.txt
   $ ./findGreater Hello file1.txt file2.txt
   file1.txt file2.txt
   $ ./findGreater thought file1.txt file2.txt
   file1.txt
   $ ./findGreater thought file2.txt file2.txt
   file2.txt file2.txt
   ```

Testing tools
Note: the scripts you write in the following questions will be useful every time you write a program. Be sure to complete them! In this course, you will be responsible for your own testing. As you fix bugs and refine your code, you will very often need to rerun old tests, to check that existing bugs have been fixed, and to ensure that no new bugs have

been introduced. This task is *greatly* simplified if you take the time to create a formal test suite, and build tools to automate your testing. In the following questions, you will implement such tools as Bash scripts.

4. Write a Bash script called `produceOutputs` that is invoked as follows:

   `./produceOutputs suite-file program`

   The argument `suite-file` is the name of a file containing a list of filename stems (more details below), and the argument `program` is the name of a program to be run.

   The `produceOutputs` script runs `program` on each test in the test suite and, for each test, creates a file that contains the output produced for that test.

   The file `suite-file` contains a list of stems, from which we construct the names of files containing the input of each test. Stems will not contain spaces. For example, suppose our suite file is called `suite.txt` and contains the following entries:

   ```
   test1 test2
   reallyBigTest
   ```

   Then our test suite consists of three tests. The first one (`test1`) will use the file `test1.in` to hold its input. The second one (`test2`) will use the file `test2.in` to hold its input. The last one (`reallyBigTest`) will use the file `reallyBigTest.in` to hold its input.

   A sample run of `produceOutputs` would be as follows:

   `./produceOutputs suite.txt ./myprogram`

   The script will then run `./myprogram` three times, once for each test specified in `suite.txt`:

   - The first time, it will run `./myprogram` with standard input redirected to come from `test1.in`. The results, captured from standard output, will be stored in `test1.out`.
   - The second time, it will run `./myprogram` with standard input redirected to come from `test2.in`. The results, captured from standard output, will be stored in `test2.out`.
   - The third time, it will run `./myprogram` with standard input redirected to come from `reallyBigTest.in`. The results, captured from standard output, will be stored in `reallyBigTest.out`.

   You can get most of the marks for this question by fulfilling the above requirements. For full marks, your script must also check for the following error conditions:

   - incorrect number of command line arguments
   - missing or unreadable `.in` files (for example, the suite file contains an entry `xxx`, but `xxx.in` doesn't exist or is unreadable).

   If such an error condition arises, print an informative error message to standard error and abort the script with a nonzero exit status.

   Note on purpose of this script: This script will be useful in situations where we provide you with a binary version of a program (but not its source code) that you must implement. By creating your own test cases (`.in` files) and then using this script to produce the intended output you will have something to compare with when you implement your own solution (see the next question for how to automate the comparisons).

5. Create a Bash script called `runSuite` that is invoked as follows:

```
./runSuite suite-file program
```

The argument `suite-file` is the name of a file containing a list of filename stems (more details below), and the argument `program` is the name of the program to be run.

In summary, the `runSuite` script runs `program` on each test in the test suite (as specified by `suite-file`) and reports on any tests whose output does not match the expected output.

The file `suite-file` contains a list of stems, from which we construct the names of files containing the input and expected output of each test. Stems will not contain spaces. For example, suppose our suite file is called `suite.txt` and contains the following entries:

```
test1 test2
reallyBigTest
```

Then our test suite consists of three tests. The first one (`test1`) will use the file `test1.in` to hold its input, and `test1.out` to hold its expected output. The second one (`test2`) will use the file `test2.in` to hold its input, and `test2.out` to hold its expected output. The last one (`reallyBigTest`) will use the file `reallyBigTest.in` to hold its input, and `reallyBigTest.out` to hold its expected output.

A sample run of `runSuite` would be as follows:

```
./runSuite suite.txt ./myprogram
```

The script will then run `./myprogram` three times, once for each test specified in `suite.txt`:

- The first time, it will run `./myprogram` with standard input redirected to come from `test1.in`. The results, captured from standard output, will be compared with `test1.out`.
- The second time, it will run `./myprogram` with standard input redirected to come from `test2.in`. The results, captured from standard output, will be compared with `test2.out`.
- The third time, it will run `./myprogram` with standard input redirected to come from `reallyBigTest.in`. The results, captured from standard output, will be compared with `reallyBigTest.out`.

If the output of a given test case differs from the expected output, print the following to standard output (assuming test `test2` failed):

```
Test failed: test2
Input:
(contents of test2.in)
Expected:
(contents of test2.out)
Actual:
(contents of the actual program output)
```

with the (`contents ...`) lines replaced with actual file contents, as described. **Follow these output specifications *very carefully*. You will lose a lot of marks if your output does not match them.** If you need to create temporary files, create them in `/tmp`, and use the `mktemp` command to prevent name duplications. **Also be sure to delete any temporary files you create in `/tmp`.**

You can get most of the marks for this question by fulfilling the above requirements. For full marks, your script must also check for the following error conditions:

- incorrect number of command line arguments
- missing or unreadable `.in` or `.out` files (for example, the suite file contains an entry `xxx`, but either `xxx.in` or `xxx.out` doesn't exist or is unreadable).

If such an error condition arises, print an informative error message to standard error and abort the script with a nonzero exit status.

6. In this question, you will generalize the `produceOutputs` and `runSuite` scripts that you created in problems 4 and 5. As they are currently written, these scripts only work for programs that take their input on stdin; the scripts cannot be used with programs that take parameters on the command line. For this problem, you will enhance `produceOutputs` and `runSuite` so that they can pass command line arguments to the program being executed. The interface to the scripts remains the same:

```
./produceOutputs suite.txt ./myprogram
./runSuite suite.txt ./myprogram
```

The format of the suite file remains the same. But now, for each `testname` in the suite file, there might be an optional `testname.args`. If the file `testname.args` is present, then the script (`produceOutputs` or `runSuite`) will run `myprogram` with the contents of `testname.args` passed on the command line and the contents of `testname.in` used for input on stdin. If `testname.args` is not present, then the behaviour is identical to problem 4/5: `myprogram` is run without arguments, and `testname.in` still supplies the input on stdin. The output of `runSuite` should be identical to the way it appeared in problem 5. All of the error-checking that was required in problems 4 and 5 is required here as well.

(a) Modify `produceOutputs` to handle argument files
(b) Modify `runSuite` to handle argument files

Note: To get this working should require only very small changes to your solution to problems 4 and 5.

Submission:
The following files are due at Due Date 1: `a1q1a.txt`, `a1q1b.txt`, `a1q1c.txt`, `a1q1d.txt`, `a1q1e.txt`, `a1q1f.txt`, `a1q1g.txt`, `a1q1h.txt`, `a1q1i.txt`, `a1q2a.txt`, `a1q2b.txt`, `a1q2c.txt`, `a1q2d.txt`, `a1q2e.txt`, `findGreater`.

The following files are due at Due Date 2: `produceOutputs`, `runSuite`, `produceOutputs`, `runSuite`.